

# THE **POCHOIR** STENCIL COMPILER


**Yuan Tang**<sup>‡\*</sup>, Rezaul Chowdhury<sup>\*</sup>, Bradley  
Kuszmaul<sup>\*</sup>, CK Luk<sup>†</sup>, Charles Leiserson<sup>\*</sup>

<sup>‡</sup>*Fudan University*, <sup>\*</sup>*MIT CSAIL* and<sup>†</sup>*Intel*

# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.

## 1D 3-point stencil






# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.


## 1D 3-point stencil



# BACKGROUND

- A **stencil code** updates every point in a  $d$ -dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1$ ,  $t-2$ , ...,  $t-k$ , for  $T$  time steps.
- Stencils are used in iterative PDE solvers such as Jacobi, multigrid, and AMR, as well as for image processing and geometric modeling.

## 1D 3-point stencil




# EXAMPLE: 2D HEAT DIFFUSION

Let  $a[t, x, y]$  be the temperature at time  $t$  at point  $(x, y)$ .

## Heat equation

$$\frac{\partial a}{\partial t} = \alpha \left( \frac{\partial^2 a}{\partial x^2} + \frac{\partial^2 a}{\partial y^2} \right)$$




$\alpha$  is the *thermal diffusivity*.



## Update rule

$$\begin{aligned} a[t, x, y] = & a[t-1, x, y] \\ & + CX \cdot (a[t-1, x+1, y] - 2 \cdot a[t-1, x, y] + a[t-1, x-1, y]) \\ & + CY \cdot (a[t-1, x, y+1] - 2 \cdot a[t-1, x, y] + a[t-1, x, y-1]) \end{aligned}$$


## 2D 5-point stencil



# CLASSIC LOOPING IMPLEMENTATION

## Implementation tricks

- Reuse storage for even and odd time steps.
- Keep a *halo* of *ghost cells* around the array with boundary values.




```
for (t = 1; t <= T; ++t) {  
  for (x = 0; x < X; ++x) {  
    for (y = 0; y < Y; ++y) { // do stencil kernel  
      a[t%2, x, y]  
        = a[(t-1)%2, x, y]  
          + CX*(a[(t-1)%2, x+1, y] - 2.0*a[(t-1)%2, x, y]  
                + a[(t-1)%2, x-1, y])  
          + CY*(a[(t-1)%2, x, y+1] - 2.0*a[(t-1)%2, x, y]  
                + a[(t-1)%2, x, y-1]);  
    } } } }
```

Conventional optimization: *loop tiling*.


# CACHE INEFFICIENCY IN LOOPING

## Example: 1D 3-point stencil



**Issue:** Looping is memory intensive and uses caches poorly. Assuming data-set size  $N$ , cache-block size  $\mathcal{B}$ , and cache size  $\mathcal{M} < N$ , the number of cache misses for  $T$  time steps is  $\Theta(NT/\mathcal{B})$ .


# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious* [FLPR99] techniques, *trapezoidal decompositions* [FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.



# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious* [FLPR99] techniques, *trapezoidal decompositions* [FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.

# ALGORITHM BEHIND POCHOIR

```
void TRAP(u, ta,tb, xa,xb, dxa,dxb, ya,yb, dya,dyb) {
     $\Delta t = tb - ta$ 
     $\Delta x = \max\{xb-xa, (xb + dxb \Delta t) - (xa + dxa \Delta t)\}$ 
     $\Delta y = \max\{yb-ya, (yb + dyb \Delta t) - (ya + dya \Delta t)\}$ 
    k = 0 // Try hyper space cut
    if ( $\Delta x \geq 2\sigma\Delta t$ ) // cut x dimension
        Trisect the zoid with x-cuts, k += 1
    if ( $\Delta y \geq 2\sigma\Delta t$ ) // cut y dimension
        Trisect the zoid with y-cuts, k += 1
    if (k > 0)
        Assign dependency levels 0, 1, ..., k to subzoids
        for i = 0 to k
            parallel for all subzoids (ta,tb, xa',xb', dxa',dxb', ya',yb', dya', dyb') with
            dependency level i
                TRAP(ta,tb, xa',xb', dxa',dxb', ya',yb', dya',dyb')
    elseif  $\Delta t > 1$  //time cut: recursively walk the lower and then upper zoids
        TRAP(ta,ta+ $\Delta t/2$ , xa,xb, dxa,dxb, ya,yb, dya,dyb)
        TRAP(ta+  $\Delta t/2$ ,tb, xa+dxa  $\Delta t/2$ ,xb+dxb  $\Delta t/2$ , dxa,dxb, ya+dya  $\Delta t/2$ ,yb+dyb  $\Delta t/2$ ,
        dya,dyb)
    else // base case
        for t = ta to tb-1
            for x = xa to xb-1
                for y = ya to yb-1
                     $u((t+1) \bmod 2, x, y) = u(t \bmod 2, x, y) + CX * (u(t \bmod 2, (x-1) \bmod X,$ 
                     $y) + u(t \bmod 2, (x+1) \bmod X, y) - 2u(t \bmod 2, x, y)) + CY * (u(t \bmod 2, x, (y-1) \bmod Y)$ 
                     $+ u(t \bmod 2, x, (y+1) \bmod Y) - 2u(t \bmod 2, x, y))$ 
                    xa += dxa
                    xb += dxb
                    yb += dya
                    yb += dyb
}
```

# ALGORITHM BEHIND POCHOIR

Difficult to write!!


```
void TRAP(u, ta,tb, xa,xb, dxa,dxb, ya,yb, dya,dyb) {
    Δt = tb - ta
    Δx = max{xb-xa, (xb + dxb Δt) - (xa + dxa Δt)}
    Δy = max{yb-ya, (yb + dyb Δt) - (ya + dya Δt)}
    k = 0 // Try hyper space cut
    if (Δx ≥ 2σΔt) // cut x dimension
        Trisect the zoid with x-cuts, k += 1
    if (Δy ≥ 2σΔt) // cut y dimension
        Trisect the zoid with y-cuts, k += 1
    if (k > 0)
        Assign dependency levels 0, 1, ..., k to subzoids
        for i = 0 to k
            parallel for all subzoids (ta,tb, xa',xb', dxa',dxb', ya',yb', dya', dyb') with
            dependency level i
                TRAP(ta,tb, xa',xb', dxa',dxb', ya',yb', dya',dyb')
    elseif Δt > 1 //time cut: recursively walk the lower and then upper zoids
        TRAP(ta,ta+Δt/2, xa,xb, dxa,dxb, ya,yb, dya,dyb)
        TRAP(ta+ Δt/2,tb, xa+dxa Δt/2,xb+dxb Δt/2, dxa,dxb, ya+dya Δt/2,yb+dyb Δt/2,
        dya,dyb)
    else // base case
        for t = ta to tb-1
            for x = xa to xb-1
                for y = ya to yb-1
                    u((t+1) mod 2, x, y) = u(t mod 2, x, y) + CX * (u(t mod 2, (x-1) mod X,
                    y) + u(t mod 2, (x+1) mod X, y) - 2u(t mod 2, x, y)) + CY * (u(t mod 2, x, (y-1) mod Y)
                    + u(t mod 2, x, (y+1) mod Y) - 2u(t mod 2, x, y))
                    xa += dxa
                    xb += dxb
                    yb += dya
                    yb += dyb
}
```

# POCHOIR STENCIL COMPILER

- Domain-specific compiler programmed in Haskell that compiles a stencil language embedded in C++, a traditionally difficult language in which to embed a separately compiled domain-specific language.
- Employs a novel cache-oblivious algorithm for arbitrary  $d$ -dimensional grids which is parallelized using Intel Cilk Plus.
- Makes it straightforward to code arbitrary periodic and nonperiodic boundary conditions, including Neumann and Dirichlet conditions.
- Implements a variety of stencil-specific optimizations.
- The stencil specification can be executed and debugged without the Pochoir compiler.

# 2D HEAT EQUATION

5-point stencil on a torus




Intel C++ compiler 12.0.0 with Cilk Plus on  
12 core Intel core i7 (Nehalem)

[pochoir@csail.mit.edu](mailto:pochoir@csail.mit.edu)

# 3D WAVE EQUATION

25-point stencil on a nonperiodic domain




Intel C++ compiler 12.0.0 with Cilk Plus on  
12 core Intel core i7 (Nehalem)

[pochoir@csail.mit.edu](mailto:pochoir@csail.mit.edu)

# 3D LATTICE BOLTZMANN METHOD

19-point stencil on a nonperiodic domain



Intel C++ compiler 12.0.0 with Cilk Plus on  
12 core Intel core i7 (Nehalem)

[pochoir@csail.mit.edu](mailto:pochoir@csail.mit.edu)



# AUTO-TUNED VS POCHOIR

	<i>Berkeley Autotuner</i>	<i>Pochoir</i>
CPU	Xeon X5550	Xeon X5650
Clock	2.66GHz	2.66 GHz
cores/socket, total	4, 8	6, 12
Hyperthreading	Enabled	Disabled
L1 data cache/core	32KB	32KB
L2 cache/core	256KB	256KB
L3 cache/socket	8MB	12 MB
Peak computation	85 GFLOPS	120 GFLOPS
Compiler	icc 10.0.0	icc 12.0.0
Linux kernel		2.6.32
Threading model	Pthreads	Cilk Plus
Problem Size	<b><math>258^3 * 1</math></b>	<b><math>258^3 * 200</math></b>
3D 7-point 8 cores	2.0 GStencil/s 15.8 GFLOPS	<b>2.49 GStencil/s</b> <b>19.92 GFLOPS</b>
3D 27-point 8 cores	<b>0.95 GStencil/s</b> <b>28.5 GFLOPS</b>	0.88 GStencil/s 26.4 GFLOPS



# PUBLICATIONS

- The Pochoir Stencil Compiler, Yuan Tang, Rezaul Alam Chowdhury, Bradley C. Kuszmaul, Chi-Keung Luk, and Charles E. Leiserson, SPAA'11: The 23rd ACM Symposium on Parallelism in Algorithms and Architectures, June 4-6, San Jose, CA, USA.
- Coding Stencil Computation using the Pochoir Stencil-Specification language, Yuan Tang, Rezaul Alam Chowdhury, Chi-Keung Luk, and Charles E. Leiserson, HotPar'11: 3rd USENIX Workshop on Hot Topics in Parallelism, May 26-27, Berkeley, CA, USA.

# THANK YOU!

EMAIL [POCHOIR@CSAIL.MIT.EDU](mailto:POCHOIR@CSAIL.MIT.EDU) TO REQUEST  
A COPY OF POCHOIR COMPILER

# OUTLINE

- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **CONCLUSION**

# OUTLINE


- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **CONCLUSION**

# FUNCTIONAL SPECIFICATION

- Embedded in C++.
- Directly executable and debuggable via any native C++ tool chain.
- Supports arbitrary  $d$ -dimensional rectangular grids.
- The stencil shape can be arbitrary.
- A point at time  $t$  can depend on points at time  $t-1, t-2, \dots, t-k$ . ---- arbitrary depth
- Both periodic and nonperiodic boundary conditions can be programmed.


# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```



# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,1}, {-1,0,-1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
```




Declare a **kernel function** kern with time parameter t and spatial parameters x and y.

```
10 Pochoir_Kernel_2D(kern, t, x, y)
11   a(t,x,y) = a(t-1,x,y)
12             + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
13             + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
14 Pochoir_Kernel_End
15 for (int x = 0; x < X; ++x)
16   for (int y = 0; y < Y; ++y)
17     a(0,x,y) = rand();
18 heat.Run(T, kern);
19 for (int x = 0; x < X; ++x)
20   for (int y = 0; y < Y; ++y)
21     cout << a(T,x,y);
22 return 0;
23 }
```



# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) -
14                      + 0.125*(a(t-1,x,y+1) -
15
16  Pochoir_Kernel_End
17  for (int x = 0; x < X; ++x)
18    for (int y = 0; y < Y; ++y)
19      a(0,x,y) = rand();
20
21  heat.Run(T, kern);
22  for (int x = 0; x < X; ++x)
23    for (int y = 0; y < Y; ++y)
24      cout << a(T,x,y);
25
26  return 0;
27 }
```




Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t, x, y)$ , we specify the corresponding cell  $\{0, 0, 0\}$ , for  $a(t-1, x+1, y)$ , we specify  $\{-1, 1, 0\}$ , and so on.



# 2D HEAT EQUATION


```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) -
14                      + 0.125*(a(t-1,x,y+1) -
15
16  Pochoir_Kernel_End
17  for (int x = 0; x < X; ++x)
18    for (int y = 0; y < Y; ++y)
19      a(0,x,y) = rand();
20
21  heat.Run(T, kern);
22  for (int x = 0; x < X; ++x)
23    for (int y = 0; y < Y; ++y)
24      cout << a(T,x,y);
25
26  return 0;
27 }
```



Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, e.g., for  $a(t, x, y)$ , we specify the corresponding cell  $\{0, 0, 0\}$ , for  $a(t-1, x+1, y)$ , we specify  $\{-1, 1, 0\}$ , and so on.

# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) -
14                      + 0.125*(a(t-1,x,y+1) -
15
16  Pochoir_Kernel_End
17  for (int x = 0; x < X; ++x)
18    for (int y = 0; y < Y; ++y)
19      a(0,x,y) = rand();
20
21  heat.Run(T, kern);
22  for (int x = 0; x < X; ++x)
23    for (int y = 0; y < Y; ++y)
24      cout << a(T,x,y);
25
26  return 0;
27 }
```



Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, e.g., for  $a(t, x, y)$ , we specify the corresponding cell  $\{0, 0, 0\}$ , for  $a(t-1, x+1, y)$ , we specify  $\{-1, 1, 0\}$ , and so on.

# 2D HEAT EQUATION


```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
```

```
5   Pochoir_Shape_2D 2D_five_pt[6]
      = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
```

```
6   Pochoir_2D heat(2D_five_pt);
7   Pochoir_Array_2D(double) a(X,Y);
8   a.Register_Boundary(zero_bdry);
9   heat.Register_Array(a);
10  Pochoir_Kernel_2D(kern, t, x, y)
11    a(t,x,y) = a(t-1,x,y)
                + 0.125*a(t-1,x+1,y)
                + 0.125*(a(t-1,x,y+1) -
```

```
12  Pochoir_Kernel_End
13  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
      a(0,x,y) = rand();
```


```
14  heat.Run(T, kern);
15  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
      cout << a(T,x,y);
18  return 0;
19 }
```



Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t,x,y)$ , we specify the corresponding cell  $\{0,0,0\}$ , for  $a(t-1,x+1,y)$ , we specify  $\{-1,1,0\}$ , and so on.

# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
```



```
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```

Declare a **boundary function** zero\_bdry on the 2-dimensional Pochoir array arr indexed by time coordinate t and spatial coordinates x and y, which always returns 0.

# BOUNDARY CONDITIONS

## Nonperiodic zero boundary

```
Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
    return 0;
Pochoir_Boundary_End
```

## Periodic (toroidal) boundary

```
#define mod(r,m) (((r) % (m)) + ((r)<0)?(m):0)
Pochoir_Boundary_2D(periodic, arr, t, x, y)
    return arr.get( t,
                    mod(x, arr.size(1)),
                    mod(y, arr.size(0)) );
Pochoir_Boundary_End
```

## Cylindrical boundary

```
#define mod(r,m) (((r) % (m)) + ((r)<0)?(m):0)
Pochoir_Boundary_2D(cylinder, arr, t, x, y)
    if (x < 0) || (x >= arr.size(1))
        return 0;
    return arr.get( t, x, mod(y, arr.size(0)) );
Pochoir_Boundary_End
```



# BOUNDARY CONDITIONS (CONT.)

## Dirichlet boundary

```
Pochoir_Boundary_2D(dirichlet, arr, t, x, y)
    return 100+0.2*t;
Pochoir_Boundary_End
```

## Neumann boundary

```
Pochoir_Boundary_2D(neumann, arr, t, x, y)
int xx(x), yy(y);
    if (x<0) xx = 0;
    if (x>=arr.size(1)) xx = arr.size(1);
    if (y<0) yy = 0;
    if (y>=arr.size(0)) yy = arr.size(0);
    return arr.get(t, xx, yy);
Pochoir_Boundary_End
```

# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,1,0}, {-1,0,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```

Initialize all points of the grid at time 0 to a random value.

# 2D HEAT EQUATION

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,1,0}, {-1,0,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```

Run a stencil computation on the Pochoir object heat for T time steps using kernel function kern. The Run method can be called multiple times.

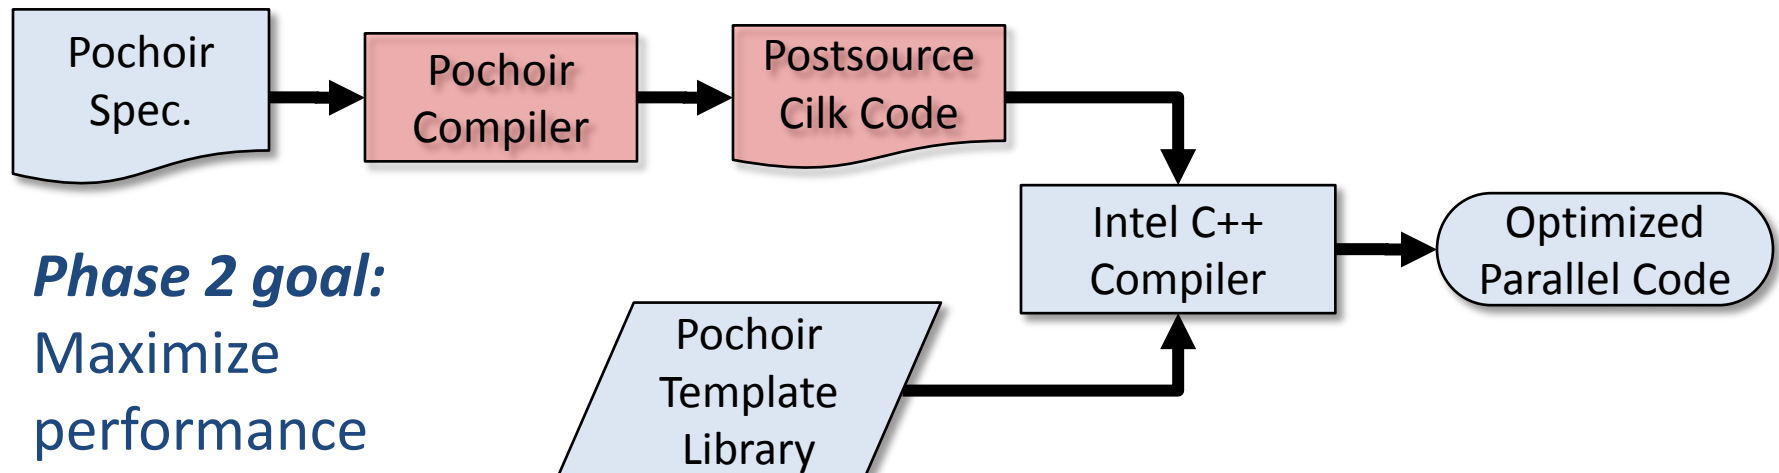



# OUTLINE

- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **CONCLUSION**

# TWO-PHASE COMPILATION STRATEGY


**Phase 1 goal:**  
Check functional  
correctness



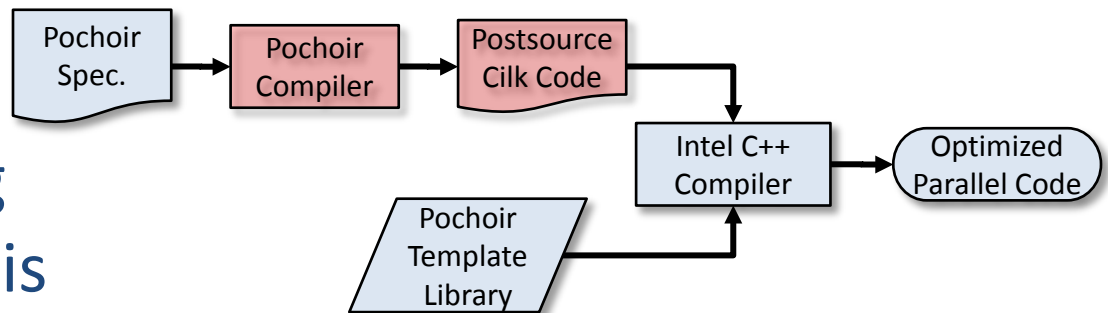
**Phase 2 goal:**  
Maximize  
performance

# POCHOIR GUARANTEE

If a stencil program compiles and runs with the Pochoir template library during Phase 1,



then no errors will occur during Phase 2 when it is compiled with the Pochoir compiler or during the subsequent running of the optimized binary.




# BENEFITS OF THE POCHOIR GUARANTEE

- The Pochoir compiler can parse as much of the programmer's C++ code as it is able without worrying about parsing it all.
- If the Pochoir compiler can “understand” the code, which it can in the common case, it can perform strong optimizations.
- If the Pochoir compiler cannot “understand” the code, it can treat the code as correct uninterpreted C++ text, confident that all the syntax- and type-checking was performed during Phase 1.

# OUTLINE


- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **CONCLUSION**

# CACHE-OBLIVIOUS STENCILS




Based on divide-and-conquer *cache-oblivious* [FLPR99] techniques, *trapezoidal decompositions* [FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.

# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious* [FLPR99] techniques, *trapezoidal decompositions* [FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.


# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious*[FLPR99] techniques, *trapezoidal decompositions*[FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.




# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious*[FLPR99] techniques, *trapezoidal decompositions*[FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.


# CACHE-OBLIVIOUS STENCILS



Based on divide-and-conquer *cache-oblivious*[FLPR99] techniques, *trapezoidal decompositions*[FS05], are asymptotically efficient, achieving  $\Theta(NT/\mathcal{MB})$  cache misses.


*How to parallelize the algorithm?*

# PARALLEL SPACE CUTS



A *parallel space cut* [FS07] produces two *black* trapezoids that can be executed in parallel and one *gray* trapezoid that must execute in series with the black trapezoids.


# PARALLEL SPACE CUTS



A *parallel space cut* [FS07] produces two *black* trapezoids that can be executed in parallel and one *gray* trapezoid that must execute in series with the black trapezoids.

# PARALLEL SPACE CUTS

*How to extend to arbitrary  $d$ -dimensional spatial-time grid?*  
*inverted trapezoid*




A *parallel space cut* [FS07] produces two *black* trapezoids that can be executed in parallel and one *gray* trapezoid that must execute in series with the black trapezoids.

# SERIAL SPACE CUTTING


```
void walk(u, t0,t1, x0,x1,dx0,dx1, y0,y1,dy0,dy1, z0,z1,dz0,dz1) {
    int dt = t1 - t0, dx = max((x1-x0), ((x1+dx1*dt)-(x0+dx0*dt))),
        dy = max((y1-y0), ((y1+dy1*dt)-(y0+dy0*dt))),
        dz = max((z1-z0), ((z1+dz1*dt)-(z0+dz0*dt)));
    if (dx>=DX_THRESH &&dx>=4*sigma_x*dt) /* cut x dimension */
        if (x1-x0 == dx) { /* cut an upright trapezoid */
            /* spawn black trapezoids */
            cilk_spawn walk(u, t0,t1, x0,x0+dx/2,dx0,-sigma_x,
                           y0,y1,dy0,dy1, z0,z1,dz0,dz1);
            walk(u, t0,t1, x0+dx/2,x1,sigma_x,dx1,
                 y0,y1,dy0,dy1, z0,z1,dz0,dz1);
            cilk_sync;
            /* spawn gray trapezoid */
            walk(u, t0,t1, x0+dx/2,x0+dx/2,-sigma_x,sigma_x,
                 y0,y1,dy0,dy1, z0,z1,dz0,dz1);
        } else { /* cut an inverted trapezoid */
            ...
        }
    }
} else if (.../* cut y dimension */) {
    ...
} else if (.../* cut z dimension */) {
    ...
} else if (.../* cut t dimension */) {
    ...
} else { /* call the base case */
    base_case(u, t0,t1, x0,x1,dx0,dx1, y0,y1,dy0,dy1, z0,z1,dz0,dz1);
}
}
```

Spatial dimensions  
are cut one at a time.

# SERIAL SPACE CUT — 2D




# SERIAL SPACE CUT — 2D







# SERIAL SPACE CUT — 2D




# SERIAL SPACE CUT — 2D



# SERIAL SPACE CUT — 2D




# SERIAL SPACE CUT — 2D



Spawn 11, 13, 31, 33  
Sync  
Spawn 12, 32  
Sync  
Spawn 21, 23  
Sync  
Spawn 22

# SERIAL SPACE CUT — 2D



Spawn 11, 13, 31, 33

Sync

Spawn 12, 32


Sync

Spawn 21, 23

Sync

Spawn 22

# SERIAL SPACE CUT — 2D



Spawn 11, 13, 31, 33

Sync

Spawn 12, 32


Sync

Spawn 21, 23


Sync

Spawn 22

# SERIAL SPACE CUT — 2D




Spawn 11, 13, 31, 33  
Sync  
Spawn 12, 32  
Sync  
Spawn 21, 23  
Sync  
Spawn 22



Spawn 11, 13, 31, 33  
Sync  
Spawn 21, 23  
Sync  
Spawn 12, 32  
Sync  
Spawn 22

# SERIAL SPACE CUT — 2D



1<sup>st</sup> cut X, 2<sup>nd</sup> cut Y

Spawn 11, 13, 31, 33

Sync


Spawn 12, 32

Sync

Spawn 21, 23

Sync

Spawn 22



1<sup>st</sup> cut Y, 2<sup>nd</sup> cut X

Spawn 11, 13, 31, 33

Sync


Spawn 21, 23

Sync

Spawn 12, 32


Sync

Spawn 22






# SERIAL SPACE CUT — 2D




Spawn 11, 13, 31, 33  
 Sync  
 Spawn 12, 32  
 Sync  
 Spawn 21, 23  
 Sync  
 Spawn 22






Spawn 11, 13, 31, 33  
 Sync  
~~Spawn 21, 23~~  
~~Sync~~  
 Spawn 12, 32  
 Sync  
 Spawn 22

# HYPER SPACE CUT




- **Space Cut:** Evaluate and assign dependency levels to as many spatial dimensions as possible. Spawn and sync subtrapezoids according to dependency levels.
- **Time Cut:** The same as Sequential Space Cut
- **Base Case:** The same as Sequential Space Cut

# HYPER SPACE CUT



Spawn 11, 13, 31, 33  
 Sync  
 Spawn 21, 23, 12, 32  
 Sync  
 Spawn 22

# SERIAL SPACE CUT VS HYPER SPACE CUT




1st cut X, 2nd cut Y    1st cut Y, 2nd cut X

Spawn  's  
 Sync  
 Spawn  's  
 Sync  
 Spawn  's  
 Sync  
 Spawn  's


Spawn  's  
 Sync  
 Spawn  's  
 Sync  
 Spawn  's  
 Sync  
 Spawn  's

pochoir@csail.mit.edu




Spawn  's  
 Sync  
 Spawn  's  
 Sync  
 Spawn  's

# HYPER SPACE CUT



**Lemma 3:** All  $(2r+1)^k$  subtrapezoids created by a hyper space cut on  $k \geq 1$  of the  $d \geq k$  spatial dimensions of a  $(d+1)$ -dimensional trapezoid can be evaluated in  $k+1$  parallel steps

# HYPER SPACE CUT




**Lemma 3:** All  $(2r+1)k$  subtrapezoids created by a hyper space cut on  $k \geq 1$  of the  $d \geq k$  spatial dimensions of a  $(d+1)$ -dimensional trapezoid can be evaluated in  $k+1$  parallel steps

**Lemma 2:** All  $(2r+1)k$  subtrapezoids created by a serial space cut on  $k \geq 1$  of the  $d \geq k$  spatial dimensions of a  $(d+1)$ -dimensional trapezoid can be evaluated in  $2^k$  parallel steps.


# SUPERIORITY OF HYPER SPACE CUTS

**Theorem.** On a  $d$ -dimensional spatial grid with all spatial dimensions roughly equal to the time dimension  $h$ , Pochoir's hyperspace-cutting algorithm achieves  $\Theta(h^{d+1-\lg(d+2)}/d)$  parallelism, while Frigo and Strumpen's original serial space-cutting algorithm achieves only  $\Theta(h^{d+1-\lg(2^{d+1})}) = O(h)$  parallelism.


Both algorithms have the same asymptotic cache complexity.



Parallelism of heat\_2D\_P




Parallelism of heat\_2D\_NP




Parallelism of 3dfd





Cache Miss Ratio of heat\_2D\_NP



Cache Miss Ratio of 3dfd

# OUTLINE

- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **RESULTS**
- **CONCLUSION**

# OPTIMIZATIONS

- Two code clones
- Unifying the handling of periodic and nonperiodic boundary conditions
- Automatic selection of traversal strategy
  - `-split-macro-shadow`
  - `-split-opt-pointer`
- Coarsening of base cases


# Two CODE CLONES

- The *slow clone* handles regions that contain boundaries and checks for out-of-range grid points.
- The *fast clone* handles the larger interior regions which require no range checking.



# Two CODE CLONES

- The *slow clone* handles regions that contain boundaries and checks for out-of-range grid points.
- The *fast clone* handles the larger interior regions which require no range checking.




During the recursive algorithm\*, the fast clone is used whenever possible.

\*The actual algorithm decomposes the grid into trapezoids, not rectangles.

# Two CODE CLONES

- The *slow clone* handles regions that contain boundaries and checks for out-of-range grid points.
- The *fast clone* handles the larger interior regions which require no range checking.




During the recursive algorithm\*, the fast clone is used whenever possible. Once the fast clone is used for a region, the fast clone is always used for its subregions.

\*The actual algorithm decomposes the grid into trapezoids, not rectangles.

# Two CODE CLONES

- The *slow clone* handles regions that contain boundaries and checks for out-of-range grid points.
- The *fast clone* handles the larger interior regions which require no range checking.




During the recursive algorithm\*, the fast clone is used whenever possible. Once the fast clone is used for a region, the fast clone is always used for its subregions.

\*The actual algorithm decomposes the grid into trapezoids, not rectangles.

# Two CODE CLONES

- The *slow clone* handles regions that contain boundaries and checks for out-of-range grid points.
- The *fast clone* handles the larger interior regions which require no range checking.




During the recursive algorithm\*, the fast clone is used whenever possible. Once the fast clone is used for a region, the fast clone is always used for its subregions.


\*The actual algorithm decomposes the grid into trapezoids, not rectangles.



# A UNIFIED ALGORITHMIC FRAMEWORK FOR BOTH PERIODIC AND NONPERIODIC BOUNDARY CONDITIONS



# A UNIFIED ALGORITHMIC FRAMEWORK FOR BOTH PERIODIC AND NONPERIODIC BOUNDARY CONDITIONS



# OUTLINE

- **FUNCTIONAL SPECIFICATION**
- **HOW THE POCHOIR SYSTEM WORKS**
- **ALGORITHMS**
- **OPTIMIZING STRATEGIES**
- **CONCLUSION**

# THE POCHOIR PROJECT

- Pochoir version 1.0 has been released under GNU GPL v3.0. Please email your name and affiliation to [pochoir@csail.mit.edu](mailto:pochoir@csail.mit.edu) to request a copy of Pochoir compiler.
- We will be interested in user feedback on usability, the Pochoir language, performance issues, feature requests, and anything else.
- We would like to collect more examples and benchmarks of stencil computations.

# FUTURE WORK

- Irregular computing domains
- Extending to Cluster (MPI)
- And much more...

# THANK YOU!

EMAIL [POCHOIR@CSAIL.MIT.EDU](mailto:POCHOIR@CSAIL.MIT.EDU) TO REQUEST  
A COPY OF POCHOIR COMPILER