

# The Pochoir Stencil Compiler Overview

Charles Leiserson, ***Yuan Tang***, Rezaul  
Chowdhury, Bradley Kuszmaul, CK Luk,  
Steven Johnson, Ekanathan Natarajan  
Mar. 13, 2012

# What's a stencil?



A stencil updates every point in a d-dimensional spatial grid at time  $t$  as a function of nearby grid points at times  $t-1, t-2, \dots, t-k$ , for  $T$  time steps.

## 2D Heat equation

$$\frac{\partial a}{\partial t} = \alpha \left( \frac{\partial^2 a}{\partial x^2} + \frac{\partial^2 a}{\partial y^2} \right)$$

$$\begin{aligned} a(t, x, y) = & a(t-1, x, y) \\ & + CX \cdot (a(t-1, x+1, y) - 2 \cdot a(t-1, x, y) + a(t-1, x-1, y)) \\ & + CY \cdot (a(t-1, x, y+1) - 2 \cdot a(t-1, x, y) + a(t-1, x, y-1)) \end{aligned}$$

## 2D 5-point stencil



# Story

- Stencils are prevailing
  - iterative PDE solvers such as Jacobi, multigrid, and AMR,
  - image processing,
  - geometric modeling.
- Highly cache-efficient stencil algorithm is known yet hard to write from case to case.
- Conventional numerical library focus on optimizing individual computation operator
- How to automate the optimization of a family of computation (such as stencil) in one framework is an open question.
  - Library?
  - DSL?
  - EDSL?
  - Compiler's pragma?
  - Autotuner?

# Roadmap

- Release 0.5 (Feb. 2011)
- Release 1.0 (Mar. 2012)
- Release 2.0 (TBD)

# Release 0.5

- Released in Feb. 2011
- Published in SPAA'11 & HotPar'11
- Simple, concise, declarative, and easily verifiable DSL embedded in C++, with Intel Cilk Plus extension.
- Arbitrary shaped, arbitrary depth stencil on arbitrary d-dimensional space-time grid, with complex boundary condition.

# Current List of Known Users

- Oscar Barenys, Univ. Politecnica of Catalonia, Spain.
- Volker Strumpen, Johannes Kepler University, Austria.
- Nicolas Pinto, MIT/Harvard
- Nicolas Vasilache, Reservoir Lab.
- Patrick S. McCormick, Los Alamos National Lab.
- Mohammed Shaheen, Max Planck Institut Informatik, Germany
- Wim Vanroose, Universiteit Antwerpen, Belgium.
- Tom Henretty, Ohio State Univ.
- Protonu Basu, Univ. of Utah.
- Shoaib Kamil, Berkeley.
- Hal Finkel, Argonne National Lab.
- Matthias Christen, Klingelbergstrass, Basel, Switzerland.
- Vinayaka Bandishti, Indian Institute of Science, Bangalore, India.
- Hans Vandierendonck, Ghent University, Belgium.

# Benchmark Suite

- Physics
  - Heat equation
  - Wave equation
  - Maxwell's equation
  - Lattice Boltzmann Method
- Computational Biology
  - RNA secondary structure prediction
  - Pairwise sequence alignment
- Computational Finance
  - American Put Stock Option Pricing
- Mechanical Engineering
  - Compressible Euler Flow
- Others
  - Conway's Game of Life
  - ...

# Release 0.5

- Algorithm:
- Performance Results:
- Specification:
- Boundary Conditions:
- Compilation strategy:
- Optimization strategies:



# Release 0.5

- Algorithm:
- Performance Results:
- Specification:
- Boundary Conditions:
- Compilation strategy:
- Optimization strategies:

# Algorithm

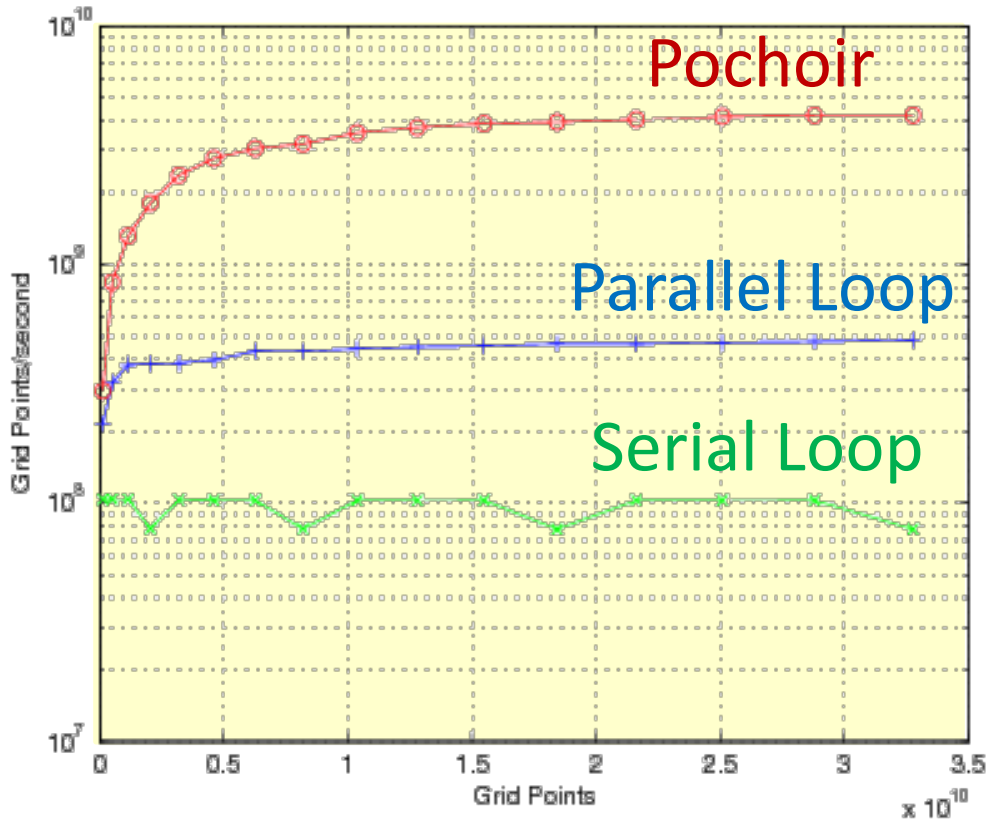
- Parallel cache-oblivious stencil algorithm
  - Based on Frigo and Strumpen's notion of trapezoidal decomposition
  - Improved parallelism by hyperspace cut strategy
  - Compared with looping implementation, reduce cache miss rate from  $\Theta(NT/\mathcal{B})$  to  $\Theta(NT/\mathcal{M}\mathcal{B})$
- C++ template meta-programming library to automatically expand for different stencils
  - Different kernels, boundaries, dimensionality, data types, etc.

# Release 0.5

- Algorithm:
- **Performance Results:**
- Specification:
- Boundary Conditions:
- Compilation strategy:
- Optimization strategies:

# 2D Heat Equation


5-point stencil on a torus



Intel C++ compiler 12.0.0 with Cilk Plus on 12  
core Intel core i7 (Nehalem)

# 3D Wave Equation


25-point stencil on a nonperiodic domain



Intel C++ compiler 12.0.0 with Cilk Plus on 12  
core Intel core i7 (Nehalem)

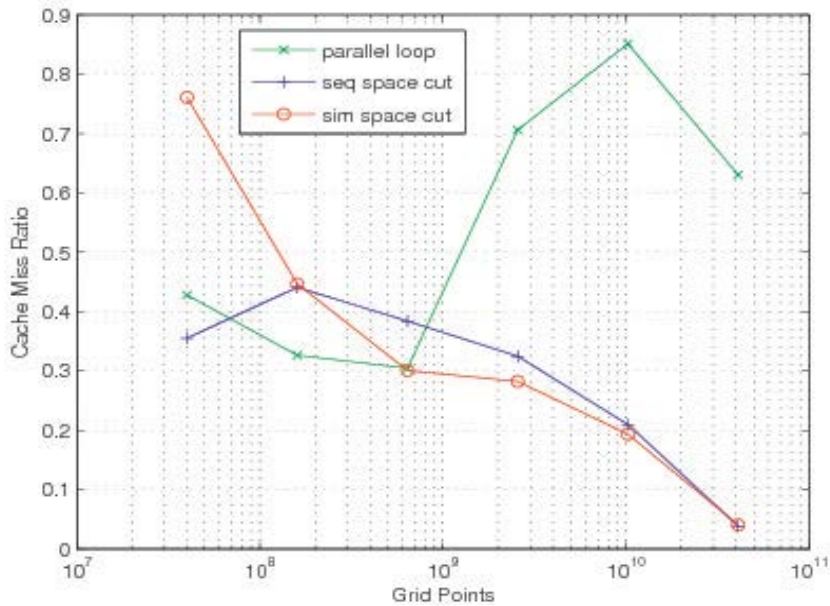
# 3D Lattice Boltzmann Method

19-point stencil on a nonperiodic domain




Intel C++ compiler 12.0.0 with Cilk Plus on 12  
core Intel core i7 (Nehalem)

# Cache miss ratio of Pochoir vs Parallel loops



Cache Miss Ratio of heat\_2D\_NP



Cache Miss Ratio of 3dfd

# Pochoir vs Autotuner

	<i>Berkeley Autotuner</i>	<i>Pochoir</i>
CPU	Xeon X5550	Xeon X5650
Clock	2.66GHz	2.66 GHz
cores/socket, total	4, 8	6, 12
Hyperthreading	Enabled	Disabled
L1 data cache/core	32KB	32KB
L2 cache/core	256KB	256KB
L3 cache/socket	8MB	12 MB
Peak computation	85 GFLOPS	120 GFLOPS
Compiler	icc 10.0.0	icc 12.0.0
Linux kernel		2.6.32
Threading model	Pthreads	Intel Cilk Plus
Problem Size	<b><math>258^3 * 1</math></b>	<b><math>258^3 * 200</math></b>
3D 7-point 8 cores	2.0 GStencil/s 15.8 GFLOPS	<b>2.49 GStencil/s</b> <b>19.92 GFLOPS</b>
3D 27-point 8 cores	<b>0.95 GStencil/s</b> <b>28.5 GFLOPS</b>	0.88 GStencil/s 26.4 GFLOPS



# Release 0.5

- Algorithm:
- Performance Results:
- **Specification:**
- Boundary Conditions:
- Compilation strategy:
- Optimization strategies:

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```

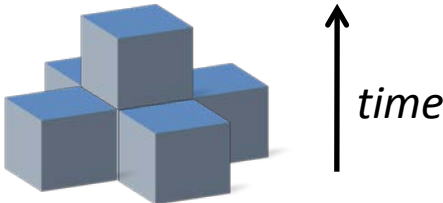


Diagram illustrating the 2D Heat Equation specification, showing a 3D grid of blue cubes representing spatial coordinates (x, y) and time (t). The time axis is indicated by an upward arrow labeled "time".

# Specification of 2D Heat Equation

```

1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End

```

```


4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,1}, {-1,0,-1}};
7   Pochoir_2D heat(2D_five_pt);

```

```

8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);

```



Declare a **kernel function** kern with time parameter  $t$  and spatial parameters  $x$  and  $y$ .

```

10 Pochoir_Kernel_2D(kern, t, x, y)
11   a(t,x,y) = a(t-1,x,y)
12             + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
13             + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
14 Pochoir_Kernel_End
15
16 for (int x = 0; x < X; ++x)
17   for (int y = 0; y < Y; ++y)
18     a(0,x,y) = rand();
19
20 heat.Run(T, kern);
21
22 for (int x = 0; x < X; ++x)
23   for (int y = 0; y < Y; ++y)
24     cout << a(T,x,y);
25
26 return 0;
27 }

```

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
```

```
5   Pochoir_Shape_2D 2D_five_pt[6]
      = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
```

```
6   Pochoir_2D heat(2D_five_pt);
```

```
7   Pochoir_Array_2D(double) a(X,Y);
```

```
8   a.Register_Boundary(zero_bdry);
```

```
9   heat.Register_Array(a);
```

```
10  Pochoir_Kernel_2D(kern, t, x, y)
```

```
11    a(t,x,y) = a(t-1,x,y)
               + 0.125*(a(t-1,x+1,y) -
               + 0.125*(a(t-1,x,y+1) -
```

```
12  Pochoir_Kernel_End
```


```
13  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
        a(0,x,y) = rand();
```

```
14  heat.Run(T, kern);
```

```
15  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
        cout << a(T,x,y);
```

```
18  return 0;
```

```
19 }
```



Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t,x,y)$ , we specify the corresponding cell  $\{0,0,0\}$ , for  $a(t-1,x+1,y)$ , we specify  $\{-1,1,0\}$ , and so on.

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
```

```
5   Pochoir_Shape_2D 2D_five_pt[6]
      = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
```

```
6   Pochoir_2D heat(2D_five_pt);
```

```
7   Pochoir_Array_2D(double) a(X,Y);
```

```
8   a.Register_Boundary(zero_bdry);
```

```
9   heat.Register_Array(a);
```

```
10  Pochoir_Kernel_2D(kern, t, x, y)
```

```
11  a(t,x,y) = a(t-1,x,y)
           + 0.125*(a(t-1,x+1,y) -
           + 0.125*(a(t-1,x,y+1) -
```

```
12  Pochoir_Kernel_End
```


```
13  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
      a(0,x,y) = rand();
```

```
14  heat.Run(T, kern);
```

```
15  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
      cout << a(T,x,y);
```

```
18  return 0;
```

```
19 }
```



Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t,x,y)$ , we specify the corresponding cell  $\{0,0,0\}$ , for  $a(t-1,x+1,y)$ , we specify  $\{-1,1,0\}$ , and so on.

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
```

```
5   Pochoir_Shape_2D 2D_five_pt[6]
      = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
```

```
6   Pochoir_2D heat(2D_five_pt);
```

```
7   Pochoir_Array_2D(double) a(X,Y);
```

```
8   a.Register_Boundary(zero_bdry);
```

```
9   heat.Register_Array(a);
```

```
10  Pochoir_Kernel_2D(kern, t, x, y)
```

```
11    a(t,x,y) = a(t-1,x,y)
                + 0.125*(a(t-1,x+1,y) -
                + 0.125*(a(t-1,x,y+1) -
```

```
12  Pochoir_Kernel_End
```


```
13  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
        a(0,x,y) = rand();
```

```
14  heat.Run(T, kern);
```

```
15  for (int x = 0; x < X; ++x)
    for (int y = 0; y < Y; ++y)
        cout << a(T,x,y);
```

```
18  return 0;
```

```
19 }
```




Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t,x,y)$ , we specify the corresponding cell  $\{0,0,0\}$ , for  $a(t-1,x+1,y)$ , we specify  $\{-1,1,0\}$ , and so on.

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
```

```
5   Pochoir_Shape_2D 2D_five_pt[6]
      = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
```

```
6   Pochoir_2D heat(2D_five_pt);
7   Pochoir_Array_2D(double) a(X,Y);
8   a.Register_Boundary(zero_bdry);
9   heat.Register_Array(a);
10  Pochoir_Kernel_2D(kern, t, x, y)
11    a(t,x,y) = a(t-1,x,y)
                + 0.125*(a(t-1,x+1,y) -
                + 0.125*(a(t-1,x,y+1) -
12  Pochoir_Kernel_End
13    for (int x = 0; x < X; ++x)
14      for (int y = 0; y < Y; ++y)
15        a(0,x,y) = rand();
16  heat.Run(T, kern);
17  for (int x = 0; x < X; ++x)
18    for (int y = 0; y < Y; ++y)
19      cout << a(T,x,y);
20  return 0;
21 }
```




Declare the 2-dimensional **Pochoir shape** 2D\_five\_pt as a list of 6 cells. Each cell specifies the relative offset of indices used in the kernel function, *e.g.*, for  $a(t,x,y)$ , we specify the corresponding cell  $\{0,0,0\}$ , for  $a(t-1,x+1,y)$ , we specify  $\{-1,1,0\}$ , and so on.

# Specification of 2D Heat Equation

```

1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End

```



```

4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,0,0}, {-1,1,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }

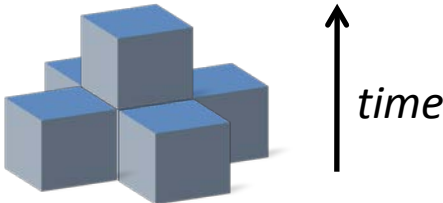
```

Declare a **boundary function** zero\_bdry on the 2-dimensional Pochoir array arr indexed by time coordinate t and spatial coordinates x and y, which always returns 0.



# Specification of 2D Heat Equation

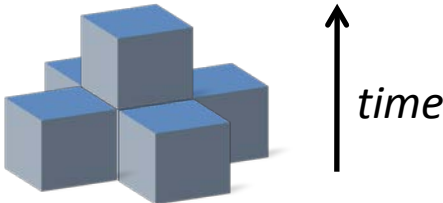
```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,1,0}, {-1,0,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```



Initialize all points of the grid at time 0 to a random value.

# Specification of 2D Heat Equation

```
1 Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
2   return 0;
3 Pochoir_Boundary_End
4 int main(void) {
5   Pochoir_Shape_2D 2D_five_pt[6]
6     = {{0,0,0}, {-1,1,0}, {-1,0,0}, {-1,-1,0}, {-1,0,-1}, {-1,0,1}};
7   Pochoir_2D heat(2D_five_pt);
8   Pochoir_Array_2D(double) a(X,Y);
9   a.Register_Boundary(zero_bdry);
10  heat.Register_Array(a);
11  Pochoir_Kernel_2D(kern, t, x, y)
12    a(t,x,y) = a(t-1,x,y)
13              + 0.125*(a(t-1,x+1,y) - 2.0*a(t-1,x,y) + a(t-1,x-1,y))
14              + 0.125*(a(t-1,x,y+1) - 2.0*a(t-1,x,y) + a(t-1,x,y-1));
15  Pochoir_Kernel_End
16  for (int x = 0; x < X; ++x)
17    for (int y = 0; y < Y; ++y)
18      a(0,x,y) = rand();
19  heat.Run(T, kern);
20  for (int x = 0; x < X; ++x)
21    for (int y = 0; y < Y; ++y)
22      cout << a(T,x,y);
23  return 0;
24 }
```



Run a stencil computation on the Pochoir object `heat` for `T` time steps using kernel function `kern`. The `Run` method can be called multiple times.

# Release 0.5

- Algorithm:
- Performance Results:
- Specification:
- **Boundary Conditions:**
- Compilation strategy:
- Optimization strategies:

# Various Boundary Conditions

## Nonperiodic zero boundary

```
Pochoir_Boundary_2D(zero_bdry, arr, t, x, y)
    return 0;
Pochoir_Boundary_End
```

## Periodic (toroidal) boundary

```
#define mod(r,m) (((r) % (m)) + ((r)<0)?(m):0)
Pochoir_Boundary_2D(periodic, arr, t, x, y)
    return arr.get( t,
                    mod(x, arr.size(1)),
                    mod(y, arr.size(0)) );
Pochoir_Boundary_End
```

## Cylindrical boundary

```
#define mod(r,m) (((r) % (m)) + ((r)<0)?(m):0)
Pochoir_Boundary_2D(cylinder, arr, t, x, y)
    if (x < 0) || (x >= arr.size(1))
        return 0;
    return arr.get( t, x, mod(y, arr.size(0)) );
Pochoir_Boundary_End
```

# Various Boundary Conditions

## Dirichlet boundary

```
Pochoir_Boundary_2D(dirichlet, arr, t, x, y)
    return 100+0.2*t;
Pochoir_Boundary_End
```

## Neumann boundary

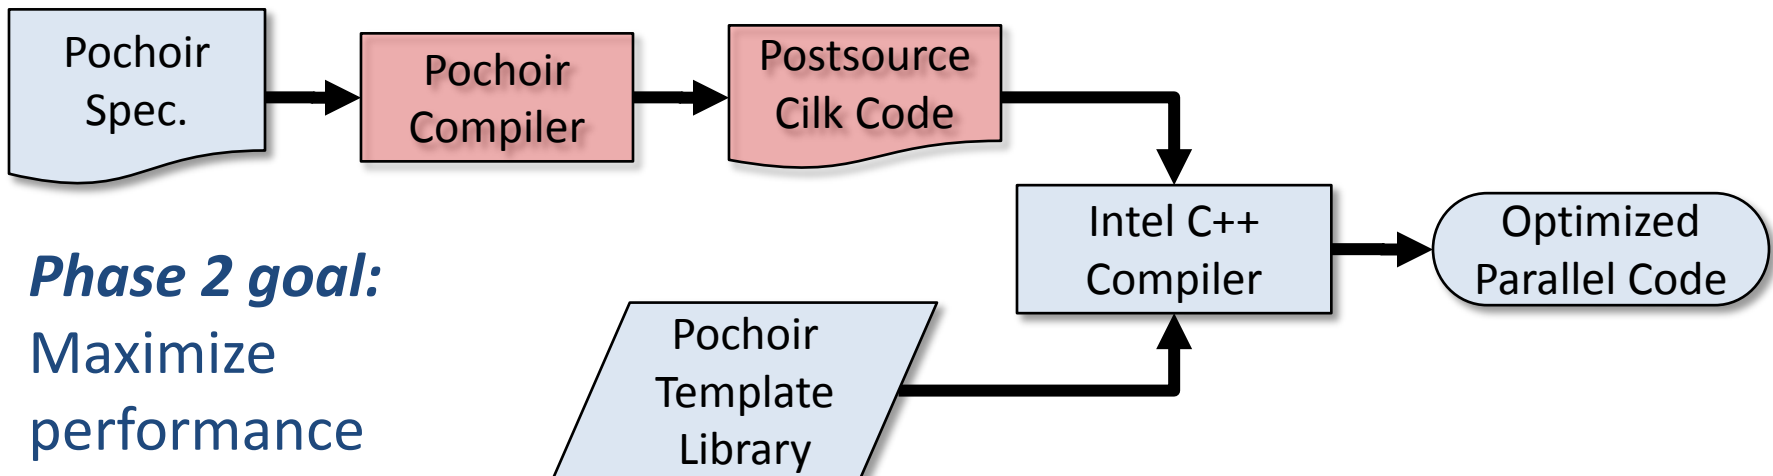

```
Pochoir_Boundary_2D(neumann, arr, t, x, y)
int xx(x), yy(y);
    if (x<0) xx = 0;
    if (x>=arr.size(1)) xx = arr.size(1);
    if (y<0) yy = 0;
    if (y>=arr.size(0)) yy = arr.size(0);
    return arr.get(t, xx, yy);
Pochoir_Boundary_End
```

# Release 0.5

- Algorithm:
- Performance Results:
- Specification:
- Boundary Conditions:
- **Compilation strategy:**
- Optimization strategies:

# Two-Phase Compilation Strategy


**Phase 1 goal:**  
Check functional  
correctness



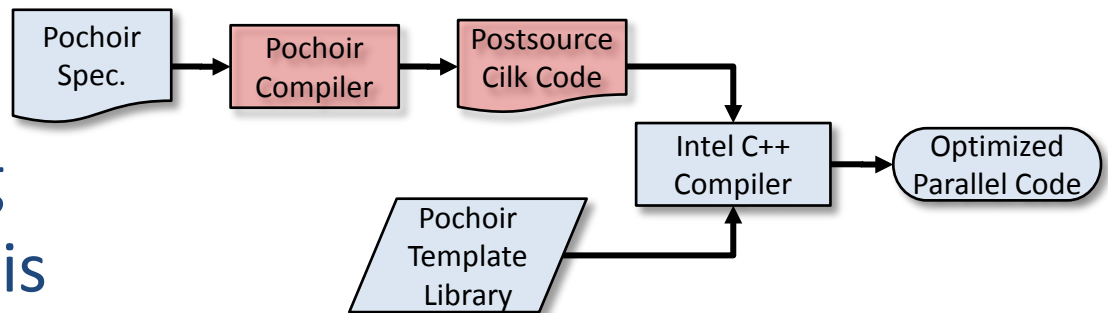
**Phase 2 goal:**  
Maximize  
performance

# Pochoir Guarantee

If a stencil program compiles and runs with the Pochoir template library during Phase 1,



then no errors will occur during Phase 2 when it is compiled with the Pochoir compiler or during the subsequent running of the optimized binary.





# Release 0.5

- Algorithm:
- Performance Results:
- Specification:
- Boundary Conditions:
- Compilation strategy:
- **Optimization strategies:**

# Optimization Strategies

- Two code clones
- Unifying the handling of periodic and nonperiodic boundary conditions
- Automatic selection of traversal strategy
  - -split-macro-shadow
  - -split-opt-pointer
- Coarsening of base cases
- Adaptive trapezoidal decomposition

# Release 1.0

- Mar. 2012
- Bug Fix
- User's feedback
- Variadic Template Support
  - Even Simpler user interface

# Release 2.0

- TBD
- Inhomogeneity
  - Macroscopic Inhomogeneity
  - Microscopic Inhomogeneity
- JIT compilation
  - Generate the computational kernels on-the-fly
- Generalized Dependency
  - From orthogonal grid to general graph

# Contributions

- Simple, concise, declarative, and easily verifiable DSL embedded in C++, with Intel Cilk Plus extension.
- Arbitrary shaped, arbitrary depth stencil on arbitrary d-dimensional space-time grid, with complex boundary condition.
- Inhomogeneous regions
  - Macroscopic inhomogeneity
  - Microscopic inhomogeneity
- JIT compiler for stencil
- Generalized dependency
  - From orthogonal grid to general graph

***Funded in medium scale by NSF***

# Pochoir Team

- Charles E. Leiserson, Project Leader
- Yuan Tang
- Rezaul Alam Chowdhury
- Bradley C. Kuszmaul
- Chi-Keung Luk
- Steven G. Johnson
- Ekanathan Palamadai Natarajan